

1. SQL

DDL: data definition language

CREATE TABLE

data types: text, float, integer, char

DML: data manipulation language

SELECT [DISTINCT] <column expression list>

FROM <single table>

[WHERE <predicate>]

[GROUP BY <column list>]

[HAVING <predicate>]

[ORDER BY <column list>]

[LIMIT <int>]

- when using group by, must have some regular expression in the select statement, or put it inside aggregate

ORDER is reading SQL logic:

FROM, WHERE, SELECT, GROUP BY, HAVING, DISTINCT, ORDER BY, LIMIT

- if a WHERE clause evaluated to null, not in output

- aggregates can't be in WHERE clause

CROSS JOINS - can't have aggregation in WHERE clause

FROM A, B

- cross product. Every row in left with every row on right

INNER JOIN

same as CROSS JOIN but need ON.

OUTER JOIN

- Even if ON predicate doesn't match, still in output

LEFT MEANS

- still if no match, left columns table are still there with nulls

RIGHT

- vice versa

Facts

- Duplicate key in Primary key not allowed by definition

Relational Algebra

no duplicates

- all operators take in relations and output different relation

Projection (π)

SELECT: select only columns specified

Selection (σ)

WHERE clause:

Page = 12, name = 'SAM'

Union (\cup)

- must be same columns

(-) set diff(\cap) intersectionJoins (\bowtie)

no specification = natural join (join on tables w/ (same names))

Rename (ρ)

cats name = drame prame \rightarrow drame (days)

PATTERN Matching

LIKE - if NO % or _ , then ' ' acts like equality
 - an underscore stands for matches single character
 - % matches any sequence of zero or more chars.

standard regex

~ . (.) represents single wildcard character
 (*) represents repetition of prev item zero or more times
 a.c = abc valid

2. Disks, Files, Buffers

Disk vs Flash (SSD)

Disk: Accessing a page

- seek time, rotational delay, transfer time
- random read much slower than sequential
- random write faster than sequential

Flash: random write faster than sequential

- faster than disk for lo-loc random I/Os
- Locality matter for both
- Disk 10x/capacity/\$ dollar


File > Page > Records

File Structures


Heap File - no particular ordering

1. Linked List implementation

- Each data page has records, free space tracker, pointers
- one header page as start



2. Page Directory implementation



- pointer to data page and free space info about that page
- faster insertion in terms of I/Os
- read header, read data, write data, write header
- fast insertion
- slow search

Sorted File - pages are ordered and records on pages sorted by keys

- implemented w/ Page Directory
- LogN search
- LogN+N insert because shift

Record Types

Fixed length - only fixed types and all same length records

Variable length - variable length. Fixed length fields before variable length fields and header has pointers to the end of the variable field

Record ID: (page #, record # on page)

Page Formats

1. Pages w/ FLR: page header to store # of records

- Packed:
- Unpacked: bitmap in header

2. Pages w/ VLR

- page footer that maintains slot directory tracking
- 1. slot count, free space pointer, entries
- 4 bytes + 4 bytes + 8 * # records
- 2. [record pointer, record length]

Heap vs Sorted

Heap: good for frequent deletes, inserts, updates. fine for frequent full scan

Sorted: good for range searches, frequent lookup

3. B+ Trees / Indexes

Index: data structure that allows fast lookup, to certain key

B+ Tree Properties

- d: order of tree. Each node except root must have $d \leq x \leq 2d$ entries assuming no delete, sorted
- Inner Node: 2d entries, 2d+1 children ptr. (tree fanout)
- The keys in the children to the left of an entry must be \leq to right 2 than.

Insertion

1. Find leaf node to insert. Add key, round to leaf
2. if overflow ($L \geq 2d$):
 - a) split into L_1 , L_2 - d in L_1 , $d+1$ in L_2
 - b) if L_1 is leaf, copy L_2 's first entry into parent, else, move
 - c) adjust pointers
3. if parent overflow, recurse on it w/ step 2.

DELETE

- just delete in leaf

Total capacity: $(2d)(2d+1)^h$ where h : # edges from RootCounting I/Os

1. Read appropriate root to leaf. one I/O per node
2. Read appropriate data page. I/O per page. (account for clustering)
3. Write data page, if modifying. if we want to write that spans multiple pages, I/O per page
4. update index page (Node)

1 I/O per linked list data page

- when searching key, there's not always a datapage associated so BEST ASE is NOT bringing I/O for data.

Fill Factor is for preventing splits on new inserts, not I/O.

Build and - build from scratch

1. sort data on key of index

2. fill leaf pages till full

3. Add ptr from parent to leaf.

If parent overflows, follow split parent.

a) keep d in L_1 , $d+1$ in L_2 b) move L_2 's first entry up.

4) adjust ptrs

Storing Records

Alt 1: leaf pages are records themselves

Alt 2: leaf pages are pointers to corresponding record

Alt 3: linked list of pointers to corresponding records

Clustering

1. Unclustered

2. Clustered

- better caching

- when multiple query, we might have that page already since sequential.

3. Sorting

Goal: sort pages in Disk

External Merge Sort

Assume N data pages in Disk, B buffer pages available

1. Load B buffer pages fully and sort that group, do that for all $\lceil N/B \rceil$ sorted runs.
2. Merge $B-1$ pages recursively

$$O(2N \cdot (1 + \lceil \log_{B-1} \lceil N/B \rceil \rceil))$$

Two Way: $2N \cdot (1 + \lceil \log_2 N \rceil)$

5. Hashing

Goal: Group same values pages in Disk

External Hashing

Assume B Buffer frames available


Divide Phase

- 1 input buffer
- $B-1$ output buffers / partitions
- group $B-1$ partitions into disk



Conquer Phase

- Start when all partitions fit in B pages
- Use fine grained hash function



Fact: 2 pass if $B(B-1)$ pages

If some of the pages/partitions are $\geq B$, first finish the pass and in the next pass recursively partition just that partition. So split that into $B-1$ as usual.

- Usually write more pages after reading \star . Round up. So, 12 data pages, 6 buffer pages, each of 5 frames get $12/5 = 3$ each. So write 15 and read 12
- Don't forget last pass of Reading / writing all the pages again fitting.

4. Buffer Management

- buffer manager works with disk space manager and RAM.
- between RAM and DISK
- $B+$ code makes read requests from Buffer Manager
- Metadata: FrameID, PageID, DirtyBit, Pin Count

Buffer Pool

Memory converted to buffer pool by partitioning into frames for pages.

1. FrameID: uniquely associated w/ memory address
2. PageID: determining which page a frame contains
3. Dirty Bit: verifying if modified
4. Pin Count: tracking number of requestors using a page

Handling Page Requests

Upon page request:

- if page exists in memory:
 1. pin count ++
 2. return address of page
- else:
 - if space:
 1. read page into it
 2. pin count = 1
 3. return page
 - else:
 REPLACEMENT POLICY!

LRU

- pro: good for repeated access
wins for random access / popular access
- cons: LRU page finding requires priority queue $O(\log N)$
sequential flooding

Clock (LRU)

- extra metadata: Ref Bit

Adding Page

1. if page exists: just return it. Don't move any clock pointers, set the reference bit
2. if doesn't exist:
 1. Move clock pointer around.
 1. if not reference bit, evict, put in, set ref bit, move clock
 2. if ref bit, unref bit, move clock

MRU

- evict most recently used unpinned page
- Pin count reduced by whoever uses the page, NOT Always BUFFER MANAGER
- file/index management code / page requestor sets dirty bit.

Conceptual Ideas

- B+ Tree:
 - API: get(key, record id)
- Buffer Manager:
 - stores pages
- Remember a B+ tree is just defined by a "page" root, so this is something we need to fetch from disk, and keep fetching the other pages that pointers to into memory.

Note 10. Parallel Query Processing

def: shared nothing - every CPU has its own disk + memory
 def: intra-operator parallelism - make one operator run as fast as possible
 def: inter-operator parallelism - make query run as fast as possible by running operators in parallel

Partitioning

1. Range Partitioning

- Each machine gets a certain range of values that it will store
- good for queries that look up on a specific key
- PRO: used in parallel sorting and parallel sort merge join

2. Hash Partitioning

- Each record is hashed and is sent to a machine matches that hash value.

3. Round Robin Partitioning

- Distribute records evenly, one by one

def: Network cost - how much data we need to send over the network

Parallel Sorting

- Range partition the table
- Perform local sort on each machine

Parallel Hashing

- Hash partition the table
- Perform local hashing on each machine

Parallel Sort Merge Join

- Range partition each table using the same ranges on the join column
- Perform local sort merge join on each machine

Parallel Grace Hash Join

- Hash partition each table using the same hash function on the join column
- Perform local grace hash join on each machine

Broadcast Join

When one big table and one small table, just send the small table to each machine and each machine does a local join.

Symmetric Hash Join (pipeline friendly)

- Build two hash tables, one for each table in the join
- When a record from R arrives, probe the hash table for S for all of the matches. When a record from S arrives, probe the hash table for R for all of the matches
- Whenever a record arrives add it to its corresponding hash-table after probing the other hash table for matches

Note 11. Transactions

def: inconsistent reads - a user reads only part of what was updated
 def: lost update - two users try to update the same record so one of the updates gets lost

def: dirty reads - one user reads an update that was never committed, aborted

def: transactions - sequence of multiple actions that should be executed as a single, logical, atomic unit. Guarantee ACID properties

Atomicity - transaction ends in two ways: either commit or aborts. Atomicity means that either all actions in the Xact happen, or none

Consistency - if DB starts out consistent, ends consistent after Xact

Isolation - Execution of each Xact is isolated from that of others

Durability - if Xact commits, effect persists, must survive failures

Ensuring Isolation Property of Transactions

Serial Schedule - run all operations of one transaction to complete before beginning the next transaction

↳ NOT efficient, want to interleave transaction actions.

Equivalence - 1. involve same transactions
 2. operations are ordered same way with same transactions
 3. Each leave the database in the same state

Serializable

we can ensure this is constraint by looking for conflicting operations

- operations are from different transactions
- both operate on same resource
- at least one is a write

• conflict serializable if conflict equivalent to a serial schedule

Dependency Graph

- one node per Xact
- Edge from T_i to T_j if:
 - an operation O_i of T_i conflicts w/ an operation O_j of T_j
 - O_i appears earlier than O_j in the schedule

• Serializable if acyclic

Deadlock

- waiting for each other to release

Avoidance

Set priority by its age: now - start-time

Pipeline Breaker

Wait-Die: if T_i has higher priority, T_i waits for T_j ; else T_i aborts

Wound-wait: if T_i has higher priority, T_j aborts; else T_i waits

Detection of Deadlock

"Waits for Graph" : one node per XACT and an edge from T_i to T_j if

- T_j holds a lock on resource X
- T_i tries to acquire a lock on resource X but T_j must release its lock on resource X before T_i can acquire its desired lock
- if cycle, shoot a XACT in cycle
- start backwards, and only draw arrow if they actually wait

2PL

can't get locks after unlock

S2PL: unlock after transaction

→ Topological Sort to get equivalent serial schedule

if two schedulers order every pair of conflicting operations the same way, then they are output equivalent (Conflict equivalent)

mainly equijoin and natural join

Notation

- $[R]$ = # pages in R
- $[S]$ = # pages in S
- $|R|$ = # records in R

Simple Nested Loop Join

- each record R, get each page of S
- $O(\text{join}(R, S))$ where R is the outer loop: $[R] + |R|[S]$

Page Oriented Nested Loop Join

- for every page in R, bring in one page S
- $[R] + [R][S]$

Block Nested Loop Join

- B-2 pages as one block of R
- $[R] + \lceil \frac{[R]}{B-2} \rceil [S]$

Index Nested Loop Join

- If we have an index on S, just look it up
- $[R] + |R| * (\text{cost to look it up matching records in S})$

Hash Join

Naïve Hash Join

- Fit R into $B-2$ pages memory and then read in each record of S and look it up in the hashtable. $[R] + [S]$ I/Os

Grace Hash Join

- 1. Repeatedly hash R and S into $B-1$ buffers so that we can get partitions of $\leq B-2$ pages
- 2. If both R and S $> B-2$ pages keep partitioning
- 3. When either R or S is small enough i.e. $\leq B-2$ pages, then load smaller one into mem and create hashtable, matching against other one. - don't care about final write

cost

- partition: $2([R] + [S])$ I/O or the partitioning cost
- matching: $[R] + [S]$

Memory Requirements

- partitioning phase divides R into $(b-1)$ runs of size $\lceil \frac{[R]}{B-1} \rceil$
- matching phase requires each $[R]/(B-1) \leq B-2$
- $R < (B-1)(B-2)$

no S constraint

- Naive join better for $R < \text{memory}$
- GJ better for $R^2 > \text{memory}$

Query Optimization Module 7

Selectivity Estimation - how much a query plan costs

Rules

- capital letter = columns
- lowercase = constants

- $X = a$: $1 / (\text{unique vals in } X)$
- $X = Y$: $1 / \max(\text{unique vals in } X, \text{unique vals in } Y)$
- $X > a$: $(\max(X) - a) / (\max(X) - \min(X) + 1)$
- cond 1 AND cond 2: $\text{selectivity}(\text{cond 1}) * \text{selectivity}(\text{cond 2})$

Selectivity of Join

- join A and B on condition A.id = B.id
- $[A][B] / \max(\text{unique vals for A.id}, \text{unique vals for B.id})$

Common Heuristics

- 1. push down projects (π) and selects (σ) as far as they go
- 2. only consider left deep plans
- 3. DON'T consider cross joining unless they are the only option

Query Optimization Contd.

Pass 1 of System R

- Single table/condition operations

- consider Fullscan - OCP I/Os
- consider Index Scan

- alt 1: Cost to reach level above leaf + num leaf read
- can stop at condition threshold, so find the leaf to start at, and keep going right at leaf level because it's sorted

- alt 2/3 indexes: (Cost to reach level above leaf) + (# leaf nodes read) + (num data pages read)

- clustered index: # data pages read is the selectivity multiplied by total # data pages
- unclustered index: I/O per each record so $\text{selectivity} * \# \text{ records}$

- example:
clustered index = $2 + 0.5[L] + 0.5[B]$
unclustered = $2 + 0.5[L] + 0.5[B]$

Evaluating Query

- Either optimal I/O query on interesting columns
- Interesting = sorted on a column used by GROUPBY or ORDERBY or used in a downstream join


Pass 2... n

- 1. At each pass i, attempt to join i tables together, each from pass i-1 and pass 1

Advance optimal plan for each set and also the optimal plan for each interesting order for each set

- * Joined Tables must be on left
- * NO CROSS JOINs are considered


- Sort Merge Join
- 1. Sort R and sort S and then iteratively check. $[R] + [S]$ at the last step.




Note 12: DB Design

When designing a database, we often use Entity-Relational Models.

def: entity: a real world object described by a set of attribute values



def: Relationship: association among two or more entities



Relationship Constraints

Thin black line - many-to-many Relationship

ex) many employees can work in many departments
many departments can have many employees

Functional Dependencies and Normalization

- $X \rightarrow Y$ means X column determines Y column in a table R.
i.e. given any two tuples in table R, if their X values are the same, then their Y values must be the same
- superkey: set of columns that determine all the columns in the table
- candidate key: minimal set of columns that determine all the columns in the table
- closure of F: F^+ : set of all FDs that are implied by F

Relation R w/ 4NF

Decomposing a Relation - Boyce Codd Normal Form

- Relation R with FDs F is in BCNF if for all $X \rightarrow A$ in F^+ , $A \subseteq X$ (called trivial form) OR X is a superkey for R.

The procedure:

If $X \rightarrow Y$ violates BCNF, R becomes $R - Y$ and XY

ex) Relation $R = \{C, S, J, D, P, Q, V\}$ key C and $F = \{JP \rightarrow C, SP \rightarrow P, J \rightarrow S\}$

- to deal w/ $SP \rightarrow P$, decompose into SDP , $CSTDQV$
- to deal w/ $J \rightarrow S$, decompose $CSTDQV$ into JS and $CJDQV$
- end up w/ SDP , JS , and $CJDQV$

Attribute Closures

A^+ Suppose you have $F = A \rightarrow B, AB \rightarrow AC, BC \rightarrow BD, DA \rightarrow C$
then A^+ closure = $\{AB, ABC, ABCD, \dots\}$

Decomposition Preservation

IFF a dependency can be applied, i.e. we have a set containing all the variables of a dependency

key constraint: each department has at most one manager but employee can be manager for 0 or more dept

Thin Arrow - 1 to many - at most 1, or 0, and 0 or more

Thick Line - participation - at least one

Thick Arrow: key + participation: at least and at most one
Department \rightarrow Manager; each dep, one manager exactly one

weak entity - entity that can be identified uniquely only w/ key of another entity (lower entity)

primary key

X^+ : set of attributes closures, set of attributes implied by X